


Certified Kubernetes Security Specialist (CKS)



Cluster setup

CIS Benchmarks & K8s Hardening

- **Standardization:** Industry-wide best practices for secure K8s configuration.
- **Scope:** Covers API Server, ETCD, Kubelet, and Controller Manager.
- **Scoring:** Provides a clear "Pass/Fail" security posture.
- **CKS Focus:** Understanding which flags (e.g., `--anonymous-auth=false`) satisfy the benchmark.



Kube-bench & Automation

- **Tooling:** Aqua Security's open-source tool to run CIS checks.
- **Remediation:** Output includes specific commands to fix vulnerabilities.
- **Continuous Compliance:** Monitoring drift in node security settings.

```
[INFO] 1 Master Node Security Configuration
[INFO] 1.1 API Server
[FAIL] 1.1.1 Ensure that the --allow-privileged argument is set to false (Scored)
[FAIL] 1.1.2 Ensure that the --anonymous-auth argument is set to false (Scored)
[PASS] 1.1.3 Ensure that the --basic-auth-file argument is not set (Scored)
[PASS] 1.1.4 Ensure that the --insecure-allow-any-token argument is not set (Scored)
[FAIL] 1.1.5 Ensure that the --kubelet-https argument is set to true (Scored)
[PASS] 1.1.6 Ensure that the --insecure-bind-address argument is not set (Scored)
[PASS] 1.1.7 Ensure that the --insecure-port argument is set to 0 (Scored)
[PASS] 1.1.8 Ensure that the --secure-port argument is not set to 0 (Scored)
[FAIL] 1.1.9 Ensure that the --profiling argument is set to false (Scored)
[FAIL] 1.1.10 Ensure that the --repair-malformed-updates argument is set to false (Scored)
[PASS] 1.1.11 Ensure that the admission control policy is not set to AlwaysAdmit (Scored)
[FAIL] 1.1.12 Ensure that the admission control policy is set to AlwaysPullImages (Scored)
[FAIL] 1.1.13 Ensure that the admission control policy is set to DenyEscalatingExec (Scored)
[FAIL] 1.1.14 Ensure that the admission control policy is set to SecurityContextDeny (Scored)
[PASS] 1.1.15 Ensure that the admission control policy is set to NamespaceLifecycle (Scored)
[FAIL] 1.1.16 Ensure that the --audit-log-path argument is set as appropriate (Scored)
[FAIL] 1.1.17 Ensure that the --audit-log-maxage argument is set to 30 or as appropriate (Scored)
[FAIL] 1.1.18 Ensure that the --audit-log-maxbackup argument is set to 10 or as appropriate (Scored)
[FAIL] 1.1.19 Ensure that the --audit-log-maxsize argument is set to 100 or as appropriate (Scored)
[PASS] 1.1.20 Ensure that the --authorization-mode argument is not set to AlwaysAllow (Scored)
[PASS] 1.1.21 Ensure that the --token-auth-file parameter is not set (Scored)
[FAIL] 1.1.22 Ensure that the --kubelet-certificate-authority argument is set as appropriate (Scored)
```

LAB: Cluster Setup & Hardening

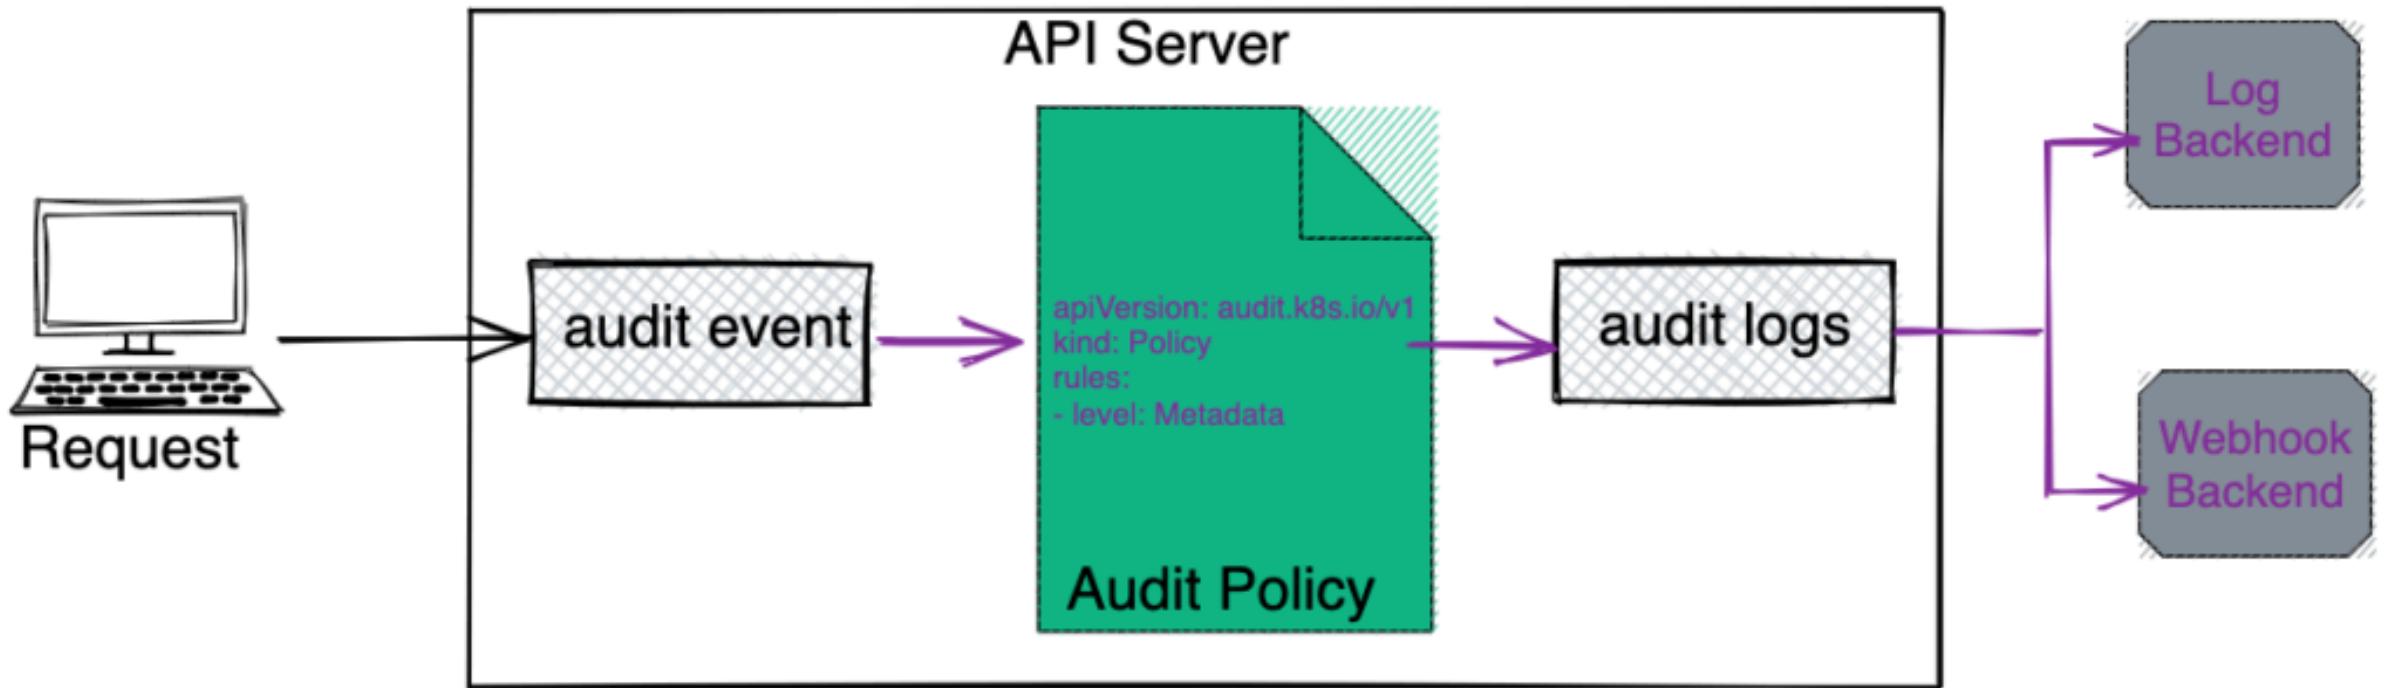
- **Exercise 1:** [Running Kube-bench as a Job](#)
- **Exercise 2:** [Manual Remediation of Control Plane](#)
- **Exercise 3:** [Automation and Continuous Compliance](#)

Auditing, Monitoring, Logging

API Server Audit Logging

- **Configuration:** Set via `--audit-policy-file` and `--audit-log-path` in `api-server` Pod.
- **⚠ Exam Tip:** You MUST know how to mount the policy file and log directory using `hostPath` volumes and `volumeMounts` within the API Server static pod manifest.

API Server Audit Workflow



Simple Audit Policy Example

```
apiVersion: audit.k8s.io/v1
kind: Policy
omitStages:
- "RequestReceived"
rules:
- level: Metadata
  resources:
  - group: ""
    resources: ["secrets"]
- level: RequestResponse
  resources:
  - group: ""
    resources: ["pods"]
  namespaces: ["production"]
```

[Official Documentation Examples](#)

Understanding Audit Stages

Each request can be recorded with an associated stage:

- **RequestReceived**: The stage for events generated as soon as the audit handler receives the request, and before it is delegated down the handler chain.
- **ResponseStarted**: Once the response headers are sent, but before the response body is sent. This stage is only generated for long-running requests (e.g. watch).
- **ResponseComplete**: The response body has been completed and no more bytes will be sent.
- **Panic**: Events generated when a panic occurred.

 **Exam Tip:** If no stages are defined in the policy, ALL stages are logged by default, which can explode the log file size.

Think of stages as a timeline of a request. For the CKS exam, you will likely use "ResponseComplete" to ensure you capture whether the request was actually authorized and successful.

Understanding Audit Levels

The defined audit levels are:

- **None**: don't log events that match this rule.
- **Metadata**: log events with metadata (requesting user, timestamp, resource, verb, etc.) but not request or response body.
- **Request**: log events with request metadata and body but not response body. This does not apply for non-resource requests.
- **RequestResponse**: log events with request metadata, request body and response body. This does not apply for non-resource requests.

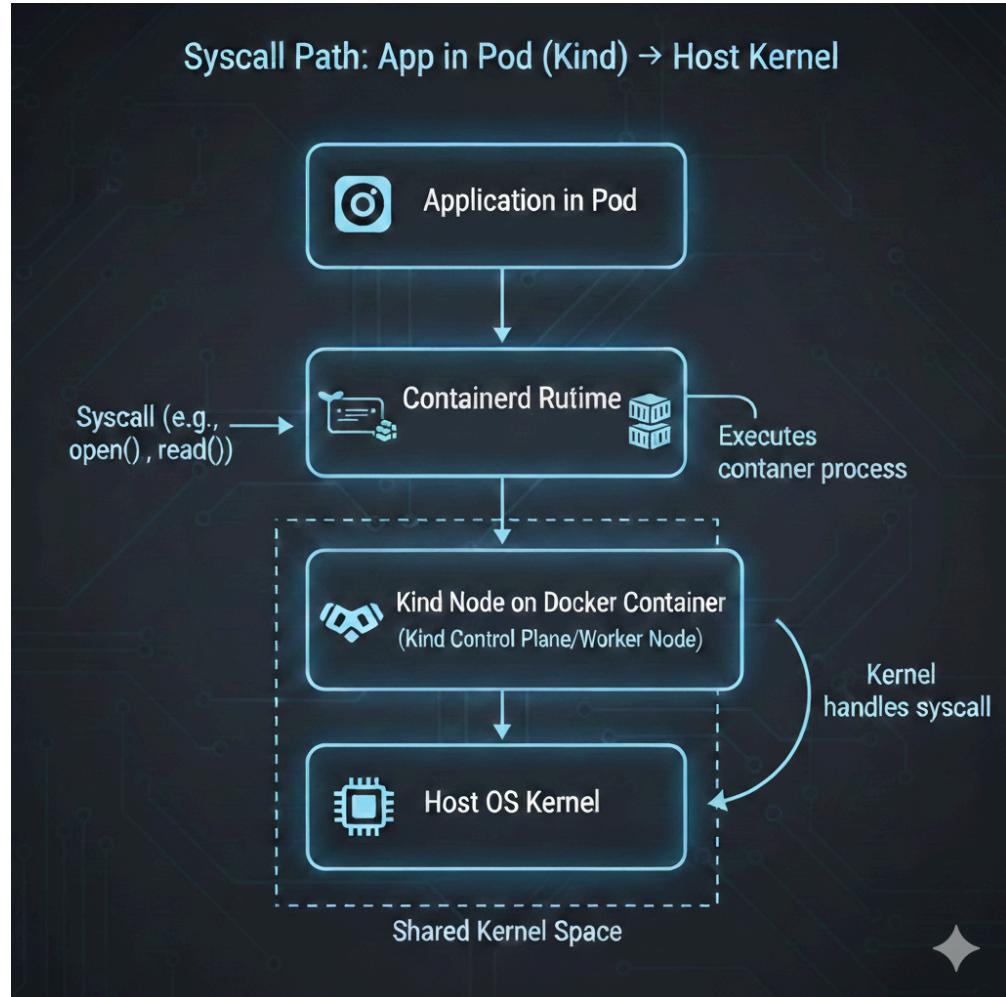
⚠ **Exam Tip:** the "RequestResponse" level is very resource-intensive for CPU and Storage. In production, we usually prefer "Metadata" for sensitive resources like Secrets to maintain performance.*

Deep Dive: Kernel & Isolation

- **Container Architecture:** Containers share the host's OS kernel through System Calls (syscalls)
- **Process Isolation:** Containers use namespaces/cgroups, but lack the hardware isolation of VMs.
- **Security Boundary:** Explains the need for syscall filtering (Seccomp).

⚠ Kind (Kubernetes in Docker)

Syscalls from a "kind" node impact the physical host kernel directly.



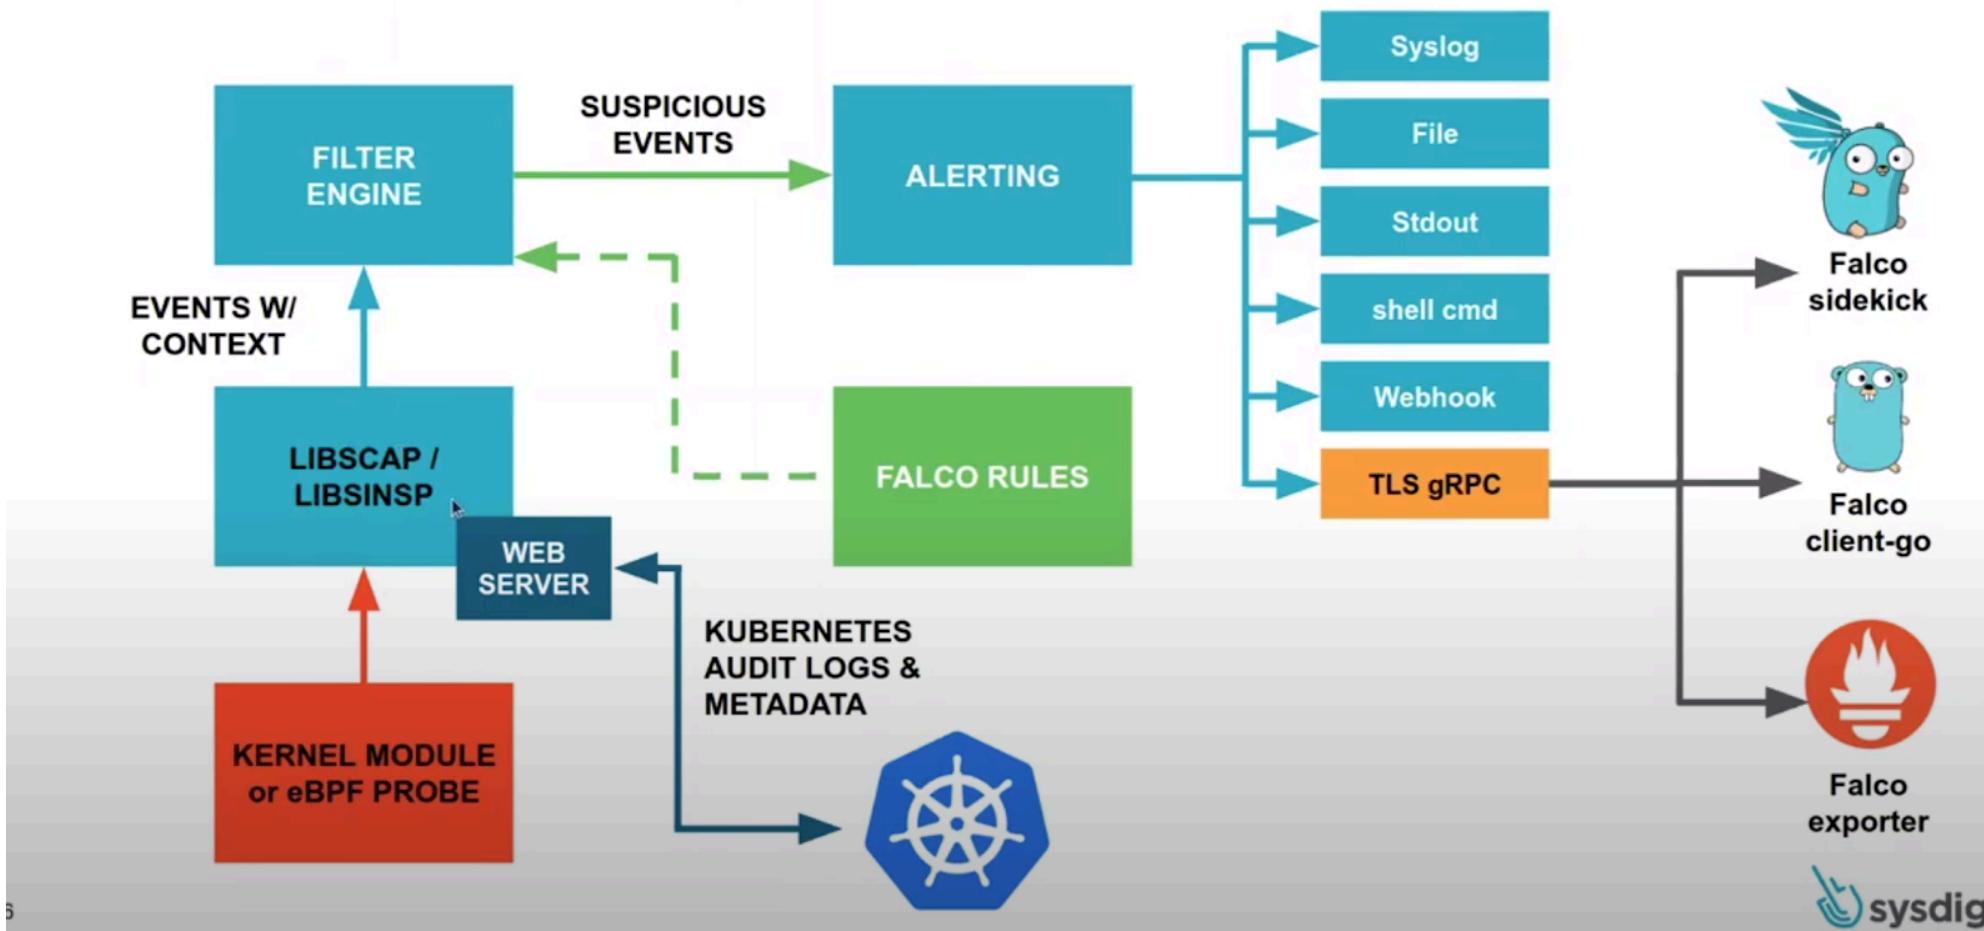
Falco: Runtime Security

- **Detection Engine:** Real-time alerting based on kernel system calls.
- **Rule Syntax:** Simple YAML format (Conditions, Outputs, Priorities).
- **Use Cases:** Detecting shell execution in pods, unauthorized file access, or outbound network changes.
- **Output:** Logs to stdout, file, gRPC, or Slack/Webhooks.

Falco is the "security monitor" of your cluster. It detects threats in real-time but does not block anything (unlike an Admission Controller)

For the CKS exam, focus on learning how to modify local rule.

Falco: Runtime Security



Falco: Rule Configuration Example

To pass the CKS, most tasks involve modifying local rule files (usually [`/etc/falco/falco_rules.local.yaml`](#)) to trigger alerts or filter false positives.

Structure of a Simple Rule

Below is a standard rule designed to detect when a shell is spawned inside a container:

```
- rule: Shell spawned in container
  desc: A shell was used inside a container
  condition: container.id != host and proc.name = bash
  output: "bash in container (user=%user.name container_id=%container.id)"
  priority: WARNING
```

Key Components for the Exam

- **Condition:** The logic used to trigger the alert. Uses field names like `fd.name` (files), `proc.name` (processes), and `container.id`.
- **Output:** The message sent to your logs. Use `%` to inject dynamic data from the event.
- **Macros & Lists:** In a real environment, you often use macros (e.g., `spawned_process`) to keep conditions readable.

Official Resources

- **Documentation:** [Falco Rules Reference](#)
- **Default Rule Examples:** [Falco GitHub Ruleset](#)

CKS Tip: If the exam asks you to "silence" an alert for a specific tool, look for the corresponding **Macro** in the configuration and add your exception there rather than deleting the whole rule.

LAB: Audit & Monitoring

- **Exercise 1:** [Configuring Audit Policies](#)
- **Exercise 2:** [Falco Runtime Alerts](#)


System Hardening

Seccomp (Secure Computing)

- **Syscall Filtering:** Limits what operations a process can request from the kernel, preventing potentially dangerous or unauthorized kernel operations.
- **Default Profiles:** `RuntimeDefault` (provided by Docker/Containerd). It blocks dangerous syscalls while allowing common operations.
- **Custom Profiles:** JSON-based filters for fine-grained syscall control, must stored on **all** Nodes in `/var/lib/kubelet/seccomp/`.

Seccomp reduces the kernel's "attack surface." If a zero-day vulnerability exists in a rarely-used syscall, Seccomp can protect the cluster by simply blocking that syscall.

Seccomp Workflow



Seccomp: Pod Application

- Use `securityContext.seccompProfile` to enforce profiles.

Example Pod with Seccomp:

```
apiVersion: v1
kind: Pod
metadata:
  name: secure-pod
spec:
  securityContext:
    seccompProfile:
      type: RuntimeDefault
  containers:
  - name: app
    image: nginx:latest
```

Seccomp profiles

- **Unconfined**: The workload runs without any seccomp restrictions.
- **RuntimeDefault**: The default Seccomp profile provided by the container runtime (Docker/Containerd).
- **Localhost**: The availability of the seccomp profile on the Node is verified by the container runtime on container creation.

Official documentation

Here is the updated content in English, integrating your specific `rsync` profile example and using the modern `securityContext` syntax.

AppArmor & SELinux

AppArmor and SELinux are Kernel-level modules that enforce Mandatory Access Control (MAC).

- restrict process capabilities—even for root
- prevent compromised containers from accessing sensitive files, unauthorized binaries, or restricted network protocols
- based on **strict security profiles**.
- **SELinux**: Label-based security. Uses contexts and policies for fine-grained system access control.
- **AppArmor**: Path-based security. Uses profiles applied to specific programs to restrict capabilities and file access.

Example: Custom AppArmor Profile (rsync)

This profile restricts the `rsync` binary to specific capabilities and file paths.

```
# File: /etc/apparmor.d/usr.bin.rsync
#include <tunables/global>

profile rsync-custom /usr/bin/rsync {
    #include <abstractions/base>
    capability net bind_service,
    network inet tcp,

    # Binaries
    /usr/bin/rsync mr,
    # Configuration files (Read-only)
    /etc/rsyncd.conf r,
    # Logs (Read/Write)
    /var/log/rsyncd.log rw,
}
```

Kubernetes Pod Implementation

Since Kubernetes v1.30, use `securityContext.appArmorProfile` to apply these profiles.

```
apiVersion: v1
kind: Pod
metadata:
  name: rsync-backup
spec:
  securityContext:
    appArmorProfile:
      type: Localhost
      # The profile "rsync-custom" must be loaded on the Node
      localhostProfile: rsync-custom
  containers:
  - name: rsync-container
    image: my-rsync-image:latest
```

Important CKS Reminders

- **Check Status:** Use `aa-status` on the node to verify the profile is in **enforce mode**.
- **Loading:** Load a new profile using: `apparmor_parser -r /etc/apparmor.d/usr.bin.rsync`
- **Localhost Type:** When using `type: Localhost`, the `localhostProfile` name must match the profile name defined in the AppArmor file exactly.
- **Deployment:** Profiles must be loaded on **all** worker nodes via `apparmor_parser` before the Pod is scheduled.

AppArmor: Key Learning Points

- Creating and loading AppArmor profiles on nodes
- Applying profiles to Kubernetes pods using `securityContext`

LAB: Kernel-Level Security

- **Exercise:** [Custom Seccomp Profiles](#)

Apparmor lab will be skipped because it is very similar to the seccomp one

- **Exercise:** [AppArmor Security Profiles](#)
- **Official Tutorial:** [Kubernetes AppArmor Documentation](#)

Container Isolation Alternatives: Sandboxed Runtimes

- **Traditional Containers:** Share host kernel directly (Docker, containerd with runc)
 - **Sandboxed Runtimes:** Add extra isolation layers
- ⚠ **CKS Context:** Understanding these alternatives helps explain why syscall filtering (Seccomp) and MAC (AppArmor/SELinux) are important for traditional container security.

- **User-space Kernel:** Intercepts syscalls before they reach the host kernel
- **Language:** Written in Go, provides a "fake" kernel interface
- **Use Case:** High security workloads that need strong isolation
- **Trade-off:** Performance overhead vs security

Security Benefit: Malicious syscalls are handled by gVisor's user-space kernel, not the host kernel.

Kata Containers

- **VM-based:** Each container runs in its own lightweight VM
- **Hardware Isolation:** True kernel isolation using hypervisor
- **Compatibility:** OCI-compliant, works with existing container tools
- **Trade-off:** Resource overhead vs maximum security

Security Benefit: Complete kernel isolation - container breaches cannot affect the host kernel directly.

gVisor in Practice

User-space Kernel Implementation

gVisor provides a **user-space kernel** written in Go that intercepts syscalls before they reach the host kernel.

--

Create RuntimeClass for gVisor

```
apiVersion: node.k8s.io/v1
kind: RuntimeClass
metadata:
  name: gvisor
handler: runsc # gVisor runtime handler
```

--

Supply Chain security

Trivy: Vulnerability Scanning

- **Scanning Capability:** Images, Filesystems, Git repos, and K8s clusters.
- **SBOM Generation:** Creating a Software Bill of Materials (CycloneDX).
- **Severity Levels:** Filter results by `CRITICAL`, `HIGH`, etc.
- **CLI Usage:** `trivy image <image_name>` is your best friend during the exam.

Speaker Notes: Trivy est l'outil de scan d'images par excellence. Insistez sur l'importance du SBOM pour la conformité de la chaîne logicielle (savoir exactement ce qui compose votre image).

ImagePolicyWebhook

- **Admission Control:** External validation before a Pod is created.
- **Mechanism:** API Server sends a JSON payload to an external endpoint.
- **Policy:** Allowed or Denied based on image tags, provenance, or scan results.
- **Configuration:** Requires a [KubeConfig](#) file for the webhook and an Admission Configuration file.

Speaker Notes: C'est le niveau ultime de la Supply Chain Security. On ne fait plus confiance au manifest, on demande à un tiers de confiance (le Webhook) de valider si l'image est autorisée.

LAB 4: Supply Chain Security

- **Exercise 1:** [Scanning with Trivy](#)
- **Exercise 2:** [Setting up ImagePolicyWebhook](#)